

Analysis of Comparison of SIMC Indexing Methods

Chang GUO, Libo ZHUO

For a range of parameter settings:

- size of tuples (determined by number of attributes)
- size of bit-slices (determined by expected number of tuples)
- size of tuple and page signatures (determined by pF)

In this project, we check and analysis comparison of parameters via following queries

- not using indexing (use x as the signature type)
- open query (?, ?, ..., ?)
- query with one solution (100001, ?, ..., ?)
- query with many solutions (?, ?, a3-001, ..., ?)
- query with multiple values (100001, ?, a3-001, ..., ?)

Each query with its implementation and shell responses as follows, it is obviously that bit slice is the fastest way to index.

```
----- Test Running... -----

----- Building Database... -----

----- Generating Data... -----
1000000,NwLRbBMqbHcdARZoWkKy,a3-000,a4-000
1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
1000002,bLDBefsaRcbYNeCDYgGX,a3-002,a4-002
1000003,XPKIOrELLnmpAPQFwKho,a3-003,a4-003
1000004,pKmcOqhNwNKueWhSQMgB,a3-004,a4-004
----- Creating Relationship... -----
----- Checking Status... -----

Global Info:
Dynamic:
  #items: tuples: 0  tsigs: 0  psigs: 0  bsigs: 0
  #pages: tuples: 1  tsigs: 1  psigs: 1  bsigs: 2
Static:
  tups  #attrs: 4  size: 42 bytes  max/page: 97
  sigs  bits/attr: 9
  tsigs size: 64 bits (8 bytes)  max/page: 511
  psigs size: 5584 bits (698 bytes)  max/page: 5
  bsigs size: 8 bits (1 bytes)  max/page: 4092
----- Inserting Data... -----
Inserting: 1000000,NwLRbBMqbHcdARZoWkKy,a3-000,a4-000
Inserting: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
```

```

Inserting: 1000002,bLDBefsaRcbYNeCDYgGX,a3-002,a4-002
Inserting: 1000003,XPKIOrELLnmpAPQFwKho,a3-003,a4-003
Inserting: 1000004,pKmcOqhNwNKueWhSQMgB,a3-004,a4-004
----- Checking Status... -----
Global Info:
Dynamic:
  #items: tuples: 5  tsigs: 5  psigs: 1  bsigs: 5584
  #pages: tuples: 1  tsigs: 1  psigs: 1  bsigs: 2
Static:
  tups  #attrs: 4  size: 42 bytes  max/page: 97
  sigs  bits/attr: 9
  tsigs size: 64 bits (8 bytes)  max/page: 511
  psigs size: 5584 bits (698 bytes)  max/page: 5
  bsigs size: 8 bits (1 bytes)  max/page: 4092
----- Showing All Data... -----
1000000,NwLRbBMqbHcdARZoWkKy,a3-000,a4-000
1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
1000002,bLDBefsaRcbYNeCDYgGX,a3-002,a4-002
1000003,XPKIOrELLnmpAPQFwKho,a3-003,a4-003
1000004,pKmcOqhNwNKueWhSQMgB,a3-004,a4-004
----- Executing Query... -----
// not using indexing (use x as the signature type)
// here we can see that it examined 5 tuples within one data page
// with no output and matched tuple
./select R 1000001,?,a3-002,? x
Query Stats:
# signatures read: 0
# sig pages read: 0
# tuples examined: 5
# data pages read: 1
# false match pages: 1

// this part we test open query
// expected output is to show all the data in the database
----- open query ---- Executing Query... -----5

./select R ?,?,?,? t
Matched Tuple: 1000000,NwLRbBMqbHcdARZoWkKy,a3-000,a4-000
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Matched Tuple: 1000002,bLDBefsaRcbYNeCDYgGX,a3-002,a4-002
Matched Tuple: 1000003,XPKIOrELLnmpAPQFwKho,a3-003,a4-003
Matched Tuple: 1000004,pKmcOqhNwNKueWhSQMgB,a3-004,a4-004
Query Stats:
# signatures read: 5

```

```

# sig pages read:    1
# tuples examined:   5
# data pages read:   1
# false match pages: 0
----- Executing Query... -----
./select R ?,?,?,? p
Matched Tuple: 1000000,NwLRbBMqbHcdARZoWkKy,a3-000,a4-000
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Matched Tuple: 1000002,bLDBefsaRcbYNeCDYgGX,a3-002,a4-002
Matched Tuple: 1000003,XPKIOrELLnmpAPQFwKho,a3-003,a4-003
Matched Tuple: 1000004,pKmcOqhNwNKueWhSQMgB,a3-004,a4-004
Query Stats:
# signatures read:    1
# sig pages read:     1
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- Executing Query... -----
./select R ?,?,?,? b
Matched Tuple: 1000000,NwLRbBMqbHcdARZoWkKy,a3-000,a4-000
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Matched Tuple: 1000002,bLDBefsaRcbYNeCDYgGX,a3-002,a4-002
Matched Tuple: 1000003,XPKIOrELLnmpAPQFwKho,a3-003,a4-003
Matched Tuple: 1000004,pKmcOqhNwNKueWhSQMgB,a3-004,a4-004
Query Stats:
# signatures read:    0
# sig pages read:     0
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- one solution ---- Executing Query... -----
./select R 1000001,?,?,? t
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Query Stats:
# signatures read:    5
# sig pages read:     1
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- Executing Query... -----
./select R 1000001,?,?,? p
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Query Stats:
# signatures read:    1

```

```

# sig pages read:    1
# tuples examined:   5
# data pages read:   1
# false match pages: 0
----- Executing Query... -----
./select R 1000001,?,?,? b
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Query Stats:
# signatures read:    0
# sig pages read:     0
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- many solutions ---- Executing Query... -----
./select R ?,?,?,? t
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Query Stats:
# signatures read:    5
# sig pages read:     1
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- Executing Query... -----
./select R ?,?,?,? p
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Query Stats:
# signatures read:    1
# sig pages read:     1
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- Executing Query... -----
./select R ?,?,?,? b
Matched Tuple: 1000001,hidDQscdxRjMowfRXSJY,a3-001,a4-001
Query Stats:
# signatures read:    0
# sig pages read:     0
# tuples examined:    5
# data pages read:    1
# false match pages: 0
----- multiple values ---- Executing Query... -----
./select R 1000001,?,a3-002,? t
Query Stats:
# signatures read:    5

```

sig pages read: 1
tuples examined: 0
data pages read: 0
false match pages: 0

----- Executing Query... -----

./select R 1000001,?,a3-002,? p

Query Stats:

signatures read: 1
sig pages read: 1
tuples examined: 5
data pages read: 1
false match pages: 1

----- Executing Query... -----

./select R 1000001,?,a3-002,? b

Query Stats:

signatures read: 0
sig pages read: 0
tuples examined: 5
data pages read: 1
false match pages: 1